



A tale of love, happiness, and asynchronous Python

Michael Dory (Spies & Assassins)

Python User Group in Princeton, November 11, 2013

Introductions

SPIES ASSASSINS

(WE'RE HIRING)

Modern Web Applications with Python



Introduction to

Tornado

O'REILLY®

*Michael Dory,
Adam Parrish
& Brendan Berg*

The story

Adam

Brendan



In short:
we needed an awesome
framework for writing APIs

Tornado!

- Scalable, non-blocking web server
- Originally developed at Friendfeed
- Open-sourced by Facebook
- Actively maintained and community-supported
- Makes developers happy

Tornado - what it is

- Scalable and fast
- Non-blocking (and/or asynchronous)
- Lightweight
- Flexible yet robust
- Pure Python goodness

Tornado - what it is **not**

- Django/Pyramid (or Rails, etc.) replacement
- Loaded with admin tools and CMS options
- Full of dependencies on libraries and practices
- Stand-alone front-end server
- Static file server*

Python framework alternatives (then)

~~django~~



~~pylons
Rapid Web Development~~

~~Bottle~~

~~weasy~~

~~CherryPy~~

Python framework alternatives (now)


~~django~~



Python framework speeds (then)

~~django~~

~~Bottle~~


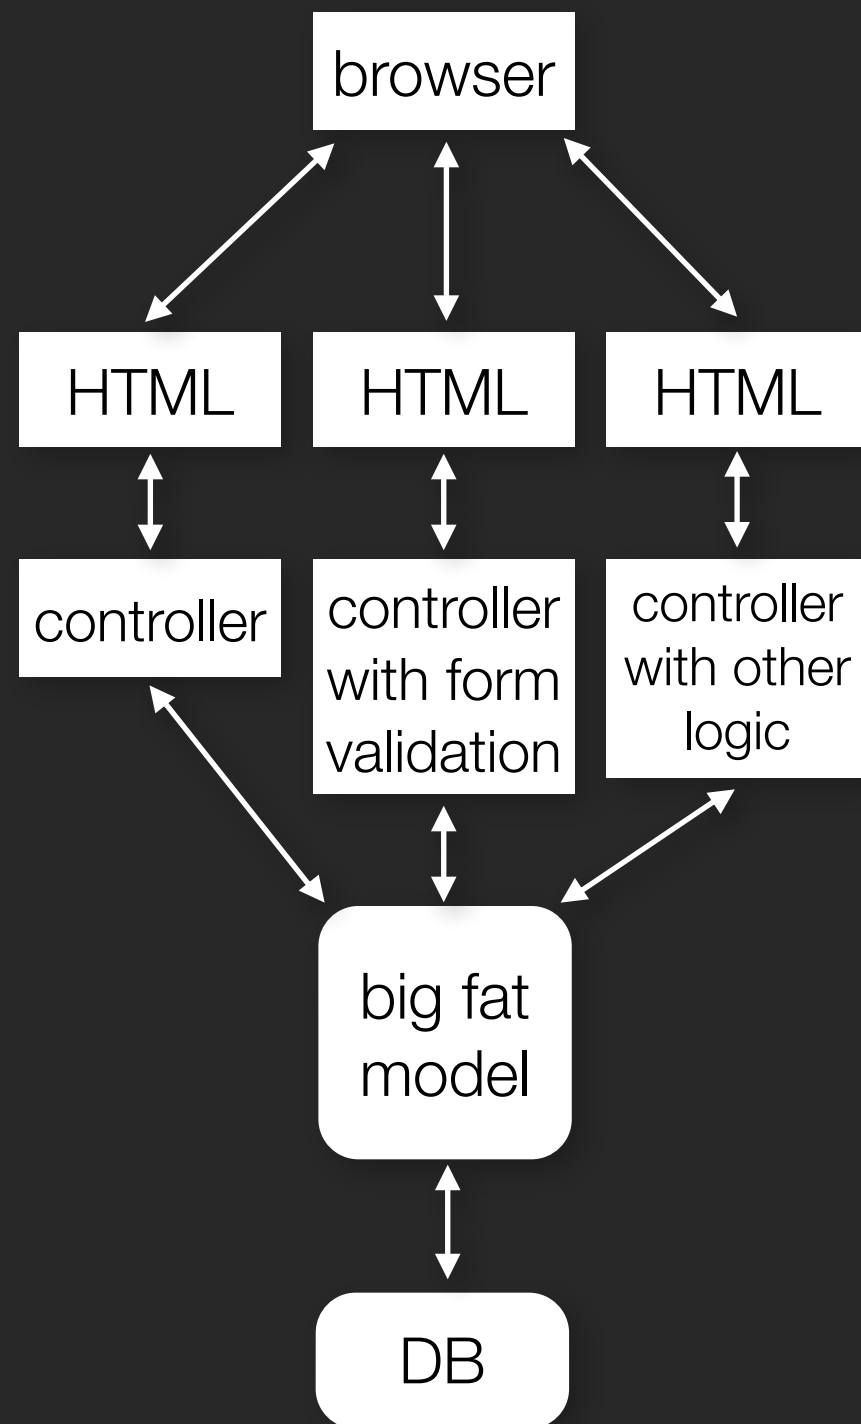


~~pylons~~
Rapid Web Development

~~CherryPy~~

Why Tornado, and not one
of the others?

Different ways of thinking



Why it might be great for you

- Quick to set up and run
- Object-oriented web handlers
- Lightweight (but powerful!) template system
- No requirements for storage frameworks

What it's perfect for

- Analytic collection (or anything real-time and write-heavy)
- Writing lightweight API's
- Social network integration
- Web applications that are I/O bound
- Apps that make use of long polling/websockets

Who else is using it?



turntable

bit.ly

10gen



hipmunk



betaworks



FINDINGS

So, that's nice.
What does it look like?

```
import os.path
import tornado.httpserver
import tornado.ioloop
import tornado.options
import tornado.web

from tornado.options import define
define("port", default=5000, help="run on the given port", type=int)

class Application(tornado.web.Application):
    def __init__(self):
        handlers = [
            (r"/([^\/]*)?", MainHandler)
        ]
        settings = dict(
            template_path=os.path.join(os.path.dirname(__file__), "templates"),
            static_path=os.path.join(os.path.dirname(__file__), "static"),
            debug=True,
        )
        tornado.web.Application.__init__(self, handlers, **settings)

class MainHandler(tornado.web.RequestHandler):
    def get(self):
        self.render(
            "main.html",
            page_title="Tornado App Example"
        )

def main():
    tornado.options.parse_command_line()
    http_server = tornado.httpserver.HTTPServer(Application())
    http_server.listen(tornado.options.options.port)
    tornado.ioloop.IOLoop.instance().start()

if __name__ == "__main__":
    main()
```

```
define("port", default=5000, help="run on the given port", type=int)

class Application(tornado.web.Application):
    def __init__(self):
        handlers = [
            (r"/([^/]+)?", MainHandler)
        ]
        settings = dict(
            template_path=os.path.join(os.path.dirname(__file__), "templates"),
            static_path=os.path.join(os.path.dirname(__file__), "static"),
            debug=True,
        )
        tornado.web.Application.__init__(self, handlers, **settings)
```



```
# main.py
class MainHandler(tornado.web.RequestHandler):
    def get(self):
        self.render(
            'main.html',
            page_title='Tornado App Example',
            page_heading='Showing a template example',
            page_items=['One', 'Two', 'Three']
        )
```



```
<!-- main.html -->
<div id="main">
    <div id="container">
        <h1>{{ page_heading }}</h1>
        <ul>
            {% for item in page_items %}
            <li>{{ item }}</li>
            {% end %}
        </ul>
    </div><!-- end container -->
</div><!-- end main -->
```

```
# main.py
class MainHandler(tornado.web.RequestHandler):
    def get(self):
        self.render(
            'main.html',
            page_title='Tornado App Example',
            page_heading='Showing a template example',
            page_items=['One', 'Two', 'Three']
        )
```



```
<!-- main.html -->
<head>
    <meta charset="utf-8">
    <meta http-equiv="X-UA-Compatible" content="IE=edge">
    <title>{{ page_title }}</title>
    <link rel="stylesheet" href="{{ static_url("css/style.css")
    }}" />
</head>
```

```
# main.py
class MainHandler(tornado.web.RequestHandler):
    def get(self):
        self.render(
            'main.html',
            page_title='Tornado App Example',
            page_heading='Showing a template example',
            page_items=['One', 'Two', 'Three']
        )
```

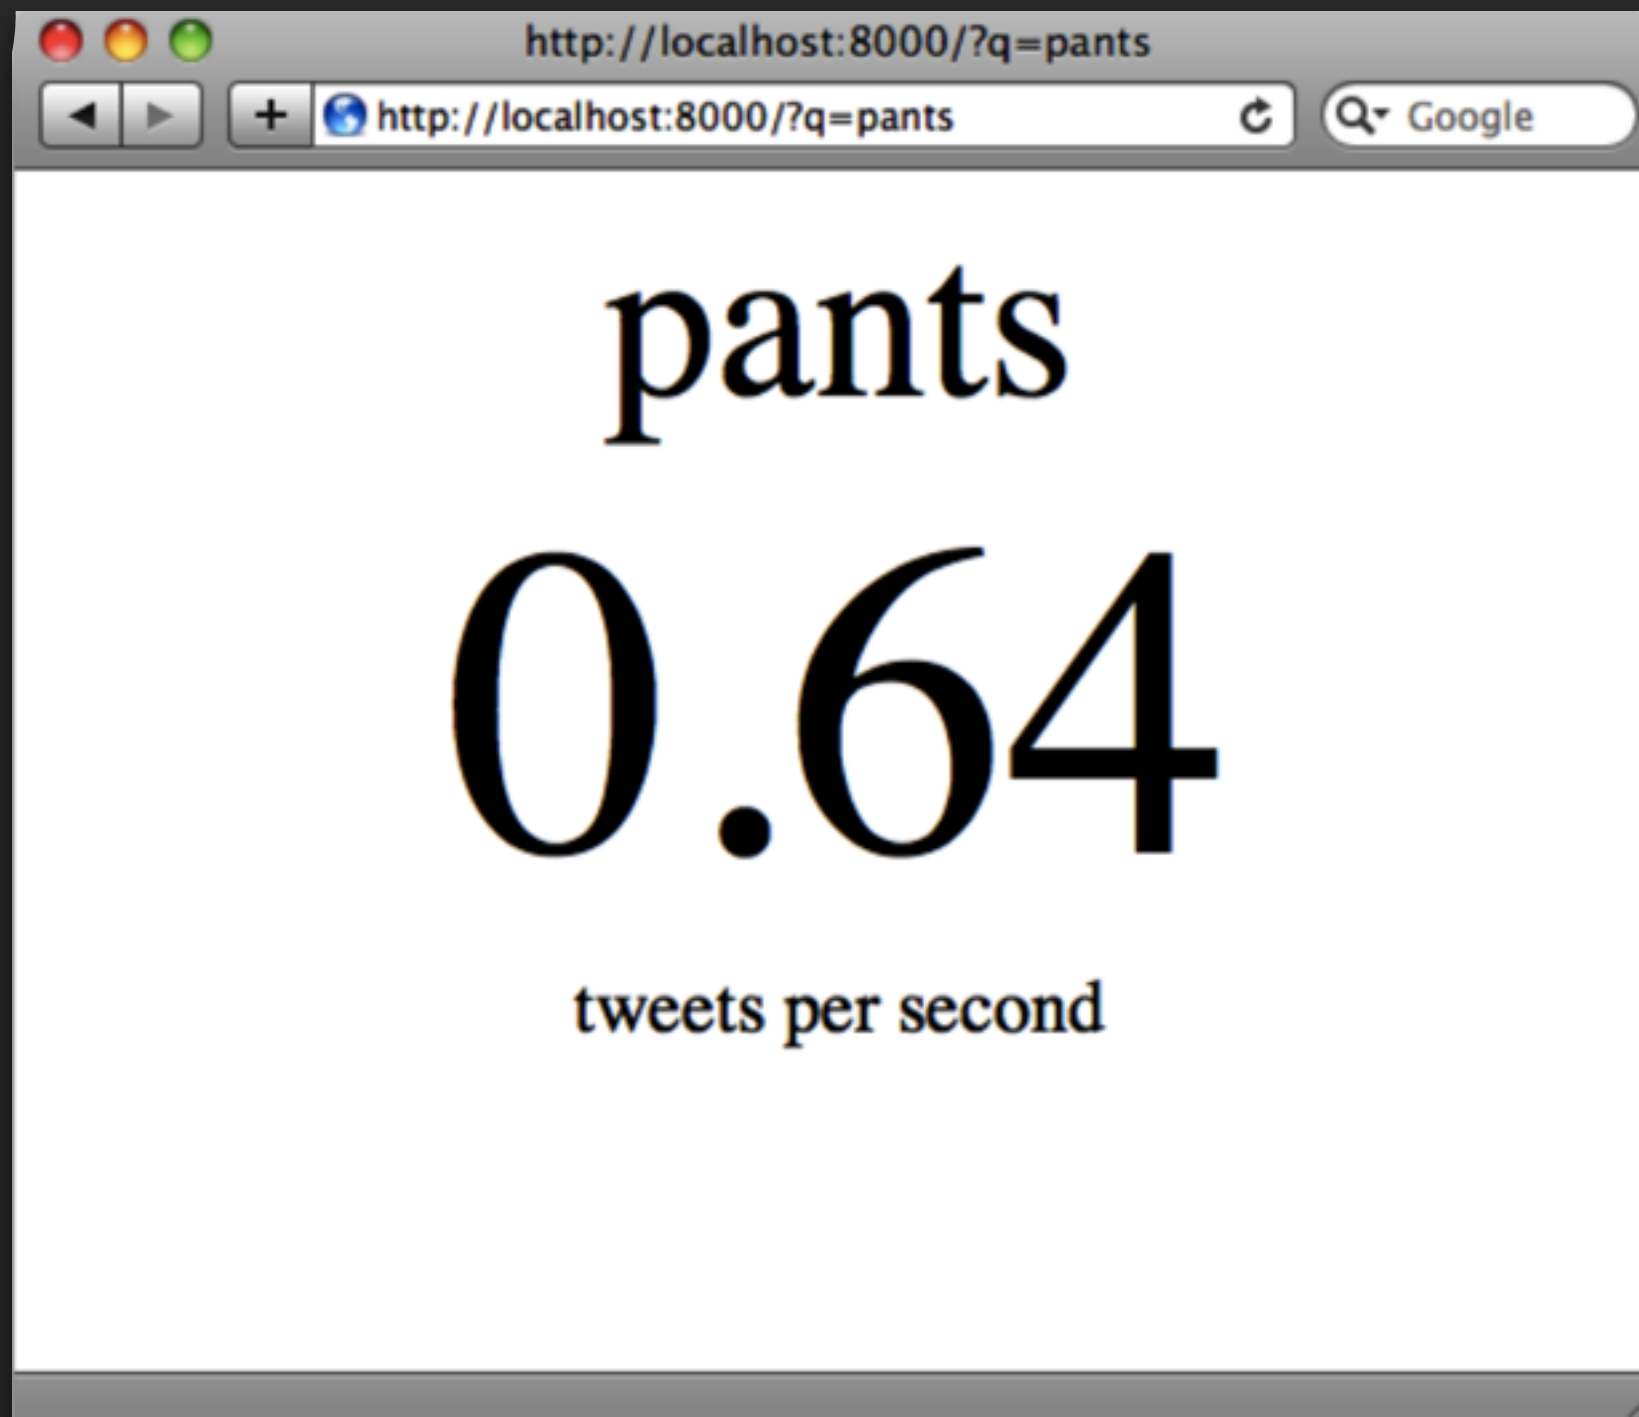


```
<!-- list-page.html -->
{% extends "main.html" %}

{% block main %}
<div id="container">
    <h1>{{ page_heading }}</h1>
    <ul>
        {% for item in page_items %}
        <li>{{ item }}</li>
        {% end %}
    </ul>
</div><!-- end container -->
{% end %}
```

What about the
asynchronous stuff?

Okay. Here's a Twitter API example.
About pants.



```

class IndexHandler(tornado.web.RequestHandler):
    def get(self):
        query = self.get_argument('q')
        client = tornado.httpclient.HTTPClient()
        response = client.fetch("http://search.twitter.com/search.json?" + \
                                urllib.urlencode(
                                    {"q": query, "result_type": "recent", "rpp": 100}))
        body = json.loads(response.body)
        result_count = len(body['results'])
        now = datetime.datetime.utcnow()
        raw_oldest_tweet_at = body['results'][-1]['created_at']
        oldest_tweet_at = datetime.datetime.strptime(raw_oldest_tweet_at,
                                                    "%a, %d %b %Y %H:%M:%S +0000")
        seconds_diff = time.mktime(now.timetuple()) - \
                       time.mktime(oldest_tweet_at.timetuple())
        tweets_per_second = float(result_count) / seconds_diff

        self.write("""
<div style="text-align: center">
    <div style="font-size: 72px">%s</div>
    <div style="font-size: 144px">%.02f</div>
    <div style="font-size: 24px">tweets per second</div>
</div>""" % (query, tweets_per_second))

```


mjd — bash — bash — 88x47
Science:~ mjd\$ siege http://localhost:8000/?q=pants -c10 -t10s

** SIEGE 2.70

** Preparing 10 concurrent users for battle.

The server is now under siege...

```
HTTP/1.1 200 0.24 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 0.48 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 0.80 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 1.07 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 0.50 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 0.69 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 0.92 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 1.14 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 1.37 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 1.61 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 1.84 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 2.01 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 2.06 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 2.30 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 2.33 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 2.33 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 2.19 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 2.25 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 1.92 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 1.99 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 1.92 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 2.15 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 3.28 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 3.27 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 3.25 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 3.18 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 3.30 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 3.63 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 3.77 secs: 178 bytes ==> /?q=pants
```

Lifting the server siege...

done.

```
Transactions: 29 hits
Availability: 100.00 %
Elapsed time: 9.07 secs
Data transferred: 0.00 MB
Response time: 1.99 secs
Transaction rate: 3.20 trans/sec
Throughput: 0.00 MB/sec
Concurrency: 6.37
Successful transactions: 29
Failed transactions: 0
Longest transaction: 3.77
Shortest transaction: 0.24
```

```

class IndexHandler(tornado.web.RequestHandler):
    @tornado.web.asynchronous
    @tornado.gen.engine
    def get(self):
        query = self.get_argument('q')
        client = tornado.httpclient.AsyncHTTPClient()
        response = yield tornado.gen.Task(client.fetch,
            "http://search.twitter.com/search.json?" + \
            urllib.urlencode(
                {"q": query, "result_type": "recent", "rpp": 100}))
        body = json.loads(response.body)
        result_count = len(body['results'])
        now = datetime.datetime.utcnow()
        raw_oldest_tweet_at = body['results'][-1]['created_at']
        oldest_tweet_at = datetime.datetime.strptime(raw_oldest_tweet_at,
            "%a, %d %b %Y %H:%M:%S +0000")
        seconds_diff = time.mktime(now.timetuple()) - \
            time.mktime(oldest_tweet_at.timetuple())
        tweets_per_second = float(result_count) / seconds_diff

        self.write("""
<div style="text-align: center">
    <div style="font-size: 72px">%s</div>
    <div style="font-size: 144px">%.02f</div>
    <div style="font-size: 24px">tweets per second</div>
</div>""" % (query, tweets_per_second))
        self.finish()

```

```
HTTP/1.1 200 0.34 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 0.22 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 0.46 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 0.24 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 0.22 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 0.24 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 0.23 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 0.21 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 0.23 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 0.24 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 0.23 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 0.23 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 0.21 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 0.23 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 0.21 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 0.22 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 0.24 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 0.22 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 1.72 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 0.23 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 0.24 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 0.23 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 0.23 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 0.22 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 0.25 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 0.22 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 0.22 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 0.20 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 0.23 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 0.22 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 0.21 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 0.21 secs: 178 bytes ==> /?q=pants
HTTP/1.1 200 0.21 secs: 178 bytes ==> /?q=pants
```

```
Lifting the server siege... done.
Transactions: 118 hits
Availability: 100.00 %
Elapsed time: 9.37 secs
Data transferred: 0.02 MB
Response time: 0.29 secs
Transaction rate: 12.59 trans/sec
Throughput: 0.00 MB/sec
Concurrency: 3.59
Successful transactions: 118
Failed transactions: 0
Longest transaction: 1.72
Shortest transaction: 0.20
```

How can I use this in
production?

NGINX



Example Nginx Configuration

```
user                nginx;
worker_processes    1;
error_log           /var/log/nginx/error.log;
pid                 /var/run/nginx.pid;

events {
    worker_connections 1024;
}

http {
    upstream tornado_pool {
        server localhost:8000;
        server localhost:8001;
        server localhost:8002;
        server localhost:8003;
    }
    server {
        listen      80;
        location / {
            proxy_set_header Host $http_host;
            proxy_redirect false;
            proxy_pass http://tornado_pool;
        }
    }
}
```

Example Supervisord Configuration

```
[unix_http_server]
file=/tmp/supervisor.sock      ; (the path to the socket file)

[supervisord]
logfile=/tmp/supervisord.log   ; (main log file;default $CWD/supervisord.log)
logfile_maxbytes=50MB          ; (max main logfile bytes b4 rotation;default 50MB)
logfile_backups=10             ; (num of main logfile rotation backups;default 10)
loglevel=info                  ; (log level;default info; others: debug,warn,trace)
pidfile=/tmp/supervisord.pid   ; (supervisord pidfile;default supervisord.pid)
nodaemon=false                 ; (start in foreground if true;default false)

[program:tornado-app]
command=/var/www/tornado-app/current/venv/bin/python app/main.py --port=8000
directory=/var/www/tornado-app/current
autostart=true
autorestart=true
```


Example Heroku Procfile

```
web: python app/main.py --port=$PORT  
# yes, that's really it.
```

Example Motor integration

```
from tornado import gen

class NewMessageHandler(tornado.web.RequestHandler):
    @tornado.web.asynchronous
    @gen.coroutine
    def post(self):
        """Insert a message."""
        msg = self.get_argument('msg')
        db = self.settings['db']

        # insert() returns a Future. Yield the Future to get the result.
        result = yield db.messages.insert({'msg': msg})

        # Success
        self.redirect('/')

```

In short:

It's quick to set up,
simple to deploy,
and easy to maintain

and makes for happy developers =)

Links!

<http://tornadoweb.org>

<https://github.com/facebook/tornado>

<http://shop.oreilly.com/product/0636920021292.do>

<https://github.com/introduction-to-tornado>

Thanks!

[@mike_dory](https://github.com/mikedory)